Standard Text CoT:
[Image + Question]
"Let me think... the red cube is left of..."
(200–500 tokens • Drifts from image)
Answer
Latent Thinking (e.g. Mull-Tokens):
[Image + Question]
[<think> + <|latent_pad|> × 20 + </think>]
(20 tokens • Zero decoding • Feed-forward width)
Answer