Narrowing to the Core Scientific Question

1. Models fail at spatial manipulation, not perception

2. Verbalising manipulation makes performance worse

3. Give model computation it does not verbalise → Latent Tokens

4. Crucial Question: Does latent computation actually happen?

Every method reports accuracy gains. Almost none audits what latents do.